

Cover Letter — Mondia (Agentic AI Engineer)

Headline — *LinkedIn note to the hiring team / email opener:*

You ask for 6+ months with Claude Code and 1+ year building agentic systems. I use Claude Code daily and I've shipped five custom MCP servers, LangGraph and CrewAI multi-agent workflows, and GraphRAG pipelines into production for a whole engineering organization.

Moamen Moustafa Basyoni Alexandria, Egypt · +20 102 786 8067 · mmoamen270@gmail.com · github.com/moamen270

Dear Hiring Team,

I'm applying for the Agentic AI Engineer role. Your requirements read like a description of my current job: agentic systems with tool use and orchestration, RAG pipelines, production Python, and hands-on Claude Code. As an R&D engineer at Andalusia I build exactly that, and it runs in production rather than in a notebook.

What I've built:

- **Five custom Model Context Protocol servers** — Azure DevOps, Figma, XMind, Playwright, and a database server — surfaced to the whole engineering organization through a custom Azure DevOps extension. This is tool-use and orchestration infrastructure, not a demo.
- **Multi-agent workflows with LangGraph and CrewAI**, including cyclic workflows and human-in-the-loop review.
- **RAG and GraphRAG pipelines** on Neo4j: a policy-retrieval graph, a code-to-business knowledge graph over a 38-module application, and a PII-masked organizational memory graph — with semantic search, contextual retrieval, and query expansion.
- **Multi-model orchestration** across OpenAI and OpenRouter with model selection, rate limiting, token budgeting, and cost optimization; LLM observability through LangSmith and LangFuse; guardrails and prompt-injection mitigation.
- Customer-facing AI products: a mail assistant for C-level users, an assessment system that generates and scores quizzes for HR and L&D, and a clinic booking assistant.

On the Python and production side: FastAPI and Flask services with Pydantic v2, async/await and dependency injection, SQLAlchemy 2.0 with Alembic over PostgreSQL, tested with pytest and pytest-asyncio, shipped through Azure DevOps CI/CD with Docker.

Two honest gaps. My Python is around two years, inside roughly three years of engineering overall, rather than the four you ask for. And while I work with Neo4j, Elasticsearch and semantic retrieval, I haven't used **pgvector** specifically — it's close to what I do and I'm getting hands-on with it, but I won't claim it.

I'd welcome a conversation.

Best regards, Moamen Basyoni