

Cover Letter — Sana Commerce (Software Engineer AI)

Headline — *LinkedIn note to the hiring team / email opener:*

I build RAG and GraphRAG pipelines, LangChain/LangGraph agents and five custom MCP servers into a live enterprise platform — and run them properly, with LangSmith/LangFuse observability, multi-model orchestration and cost controls. Alexandria-based.

Moamen Moustafa Basyoni Alexandria, Egypt · +20 102 786 8067 · mmoamen270@gmail.com · github.com/moamen270

Dear Hiring Team,

I'm applying for the Software Engineer AI role. What I'd bring is LLM engineering that already runs in production alongside real backend and system-design experience — the two halves your description asks for — and I'm based here in Alexandria.

Retrieval and agents. I've built RAG and GraphRAG pipelines on Neo4j: a policy-retrieval knowledge graph, a code-to-business graph spanning a 38-module application, and a PII-masked organizational memory graph, with semantic search, contextual retrieval and query expansion. On the agent side, multi-agent workflows with **LangChain**, **LangGraph** and **CrewAI**, plus five custom Model Context Protocol servers that give agents real tool access — Azure DevOps, Figma, XMind, Playwright and a database server — surfaced organization-wide through a custom Azure DevOps extension.

LLMOps. This is the part I think separates production work from demos. I run **LangSmith** and **LangFuse** for observability and prompt telemetry, orchestrate across **OpenAI** and **OpenRouter** with model selection, rate limiting, token budgeting and cost optimization, and manage the OpenRouter organization. Guardrails, prompt-injection mitigation, human-in-the-loop review and PII masking are all part of it.

Backend and system design. Web backends on ASP.NET Core and Python/FastAPI over PostgreSQL and SQL Server, inside a distributed architecture with Temporal Saga orchestration, RabbitMQ, service discovery and an API gateway. Cloud across **AWS** (EC2, S3) and Azure, having evaluated and chosen between them per project.

Two honest gaps. I'm at roughly three years rather than four-plus. And while I orchestrate multiple model providers, I haven't used **AWS Bedrock** specifically, and my retrieval runs on Neo4j and Elasticsearch rather than a dedicated vector store — both are adjacent to what I do daily, but I won't list them as experience I don't have.

I'd welcome a conversation.

Best regards, Moamen Basyoni